

Task 2 — Values Audit

1. Company Mission and Core Values

The proposed company for this project is a hypothetical RegTech startup named RegSphere AI Governance Solutions.

RegSphere AI Governance Solutions is an early-stage RegTech company specializing in AI governance and cross-jurisdictional compliance monitoring for financial institutions. The company focuses on helping banks manage regulatory risks arising from the use of artificial intelligence and machine learning systems in consumer finance. The company's mission is to improve trust, transparency, and accountability in AI-driven financial decision-making.

The company's core values include:

- Fairness;
- Transparency;
- Accountability;
- Regulatory adaptability;
- Responsible automation.

The company's primary strength lies in translating complex regulatory requirements into configurable governance rules and monitoring workflows. Rather than replacing legal or compliance professionals, the company aims to provide scalable monitoring infrastructure that helps institutions identify governance risks earlier and respond more consistently across jurisdictions.

At its current stage, the company is positioned as a specialized B2B RegTech provider targeting large international financial institutions. The company's long-term objective is to expand geographically by supporting multiple jurisdictions and regulatory domains, including AI governance, model risk management, and fair lending compliance.

For the purpose of this assignment, RegSphere is assumed to be a small early-stage startup with around 15 to 25 employees, including compliance specialists, data scientists, and software engineers. In the short term, it aims to serve a small number of large bank clients through enterprise subscriptions, while controlling costs by focusing on one narrow use case: AI credit governance for the US and EU. Over time, it would expand to additional jurisdictions and adjacent domains such as model risk management and AML governance.

2. Whose Perspective Does the Tool Serve?

The primary users of the proposed tool are:

- Compliance teams;
- Model risk management teams;
- Chief Compliance Officers (CCOs);
- AI governance committees within financial institutions.

However, the tool also indirectly serves:

- Regulators;
- Consumers;
- Public trust objectives.

There are clear tensions between these stakeholder groups. Financial institutions generally seek operational efficiency, lower compliance costs, and scalable automation. Regulators prioritize governance, fairness, accountability, and consumer protection. Consumers expect fair access to credit and understandable explanations for adverse decisions.

These objectives are not always perfectly aligned. For example, increasing human oversight may improve governance quality but also increase operational costs and slow down loan approval processes. Similarly, stricter fairness thresholds may reduce regulatory risk while simultaneously affecting predictive performance and profitability. The proposed tool attempts to balance these competing objectives by focusing on governance transparency and configurable monitoring rather than fully automated compliance decisions. The decision to focus on AI governance and fair lending is also commercially realistic.

Large financial institutions face increasing pressure to govern AI systems responsibly while continuing to scale digital lending operations. As regulatory scrutiny increases, institutions are likely willing to invest in governance infrastructure that reduces legal, reputational, and operational risk. The product choice is therefore driven by both ability to pay and company competence. Large international banks have strong incentives and budgets to manage AI governance risk, while RegSphere's core competence is translating regulatory requirements into configurable monitoring rules, escalation workflows, and governance documentation.

3. Is the Tool Measuring Genuine Risk or Merely Documenting Compliance?

The proposed tool is designed to measure genuine governance and fairness risk rather than simply documenting compliance activities. A key design choice reflecting this philosophy is the inclusion of fairness monitoring and model drift detection.

The tool continuously evaluates metrics such as:

- Approval rate disparity;
- False positive rate differences;
- False negative rate differences;
- Model performance drift over time.

These monitoring functions are intended to identify emerging governance risks before they become major regulatory issues. This goes beyond checklist-based compliance documentation. For example, the tool does not simply confirm that adverse action explanations exist. Instead, it evaluates whether model outputs create patterns of potentially discriminatory outcomes requiring further governance review.

Similarly, the tool includes drift monitoring even though some jurisdictions may not explicitly require a specific quantitative drift threshold. This reflects a proactive risk management philosophy rather than a purely documentation-oriented compliance approach.

At the same time, the tool intentionally does not attempt to fully automate legal interpretation or final compliance decisions. Human judgement remains necessary because AI governance and fairness assessments often involve subjective trade-offs and contextual regulatory interpretation.

4. Who Bears the Cost if the Tool Gets It Wrong?

The consequences of governance failures can affect multiple stakeholders.

False Positives

False positives occur when the tool incorrectly identifies a governance or fairness issue.

The primary costs include:

- Increased operational burden;
- Additional manual reviews;
- Delayed approvals;
- Higher compliance costs;
- Reduced operational efficiency.

These costs are primarily borne by financial institutions. However, excessive false positives may also negatively affect consumers if legitimate applications experience unnecessary delays.

False Negatives

False negatives occur when the tool fails to detect genuine governance or fairness risks. This is potentially more serious.

Possible consequences include:

- Discriminatory lending outcomes;
- Regulatory penalties;
- Consumer harm;
- Reputational damage;
- Legal liability.

In this case, consumers are likely to experience the greatest direct harm, while financial institutions face substantial regulatory and reputational consequences.

Performance Drift

Model drift creates additional risks because models may become less fair or less accurate over time as data distributions change. If drift is not detected properly, the institution may unknowingly continue operating a deteriorating model.

This can result in:

- Poor lending decisions;
- Governance failures;
- Increased bias risk;
- Reduced consumer trust.

Jurisdictional Misconfiguration

One of the most important risks in cross-border AI governance is jurisdictional misconfiguration. For example, if EU governance requirements are incorrectly configured using US regulatory assumptions, the institution may fail to meet high-risk AI obligations under the EU AI Act.

This could lead to:

- Governance failures;
- Regulatory breaches;
- Enforcement action;
- Significant reputational damage.

For this reason, the proposed tool emphasizes configurable jurisdiction-specific rule management and governance workflows.